

Training

May 16, 2022

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, accuracy_score, \
    classification_report, multilabel_confusion_matrix
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPool2D, Flatten, Dense, Dropout, \
    BatchNormalization
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

```
[2]: train = pd.read_csv('./sign_mnist_train.csv')
test = pd.read_csv('./sign_mnist_test.csv')
```

```
[3]: train.head()
```

```
[3]:   label  pixel1  pixel2  pixel3  pixel4  pixel5  pixel6  pixel7  pixel8  \
0      3     107     118     127     134     139     143     146     150
1      6     155     157     156     156     156     157     156     158
2      2     187     188     188     187     187     186     187     188
3      2     211     211     212     212     211     210     211     210
4     13     164     167     170     172     176     179     180     184

      pixel9  ...  pixel775  pixel776  pixel777  pixel778  pixel779  pixel780  \
0      153  ...      207      207      207      207      206      206
1      158  ...        69      149      128        87        94      163
2      187  ...      202      201      200      199      198      199
3      210  ...      235      234      233      231      230      226
4      185  ...        92      105      105      108      133      163

      pixel781  pixel782  pixel783  pixel784
0          206          204          203          202
1          175          103          135          149
2          198          195          194          195
```

| | | | | |
|---|-----|-----|-----|-----|
| 3 | 225 | 222 | 229 | 163 |
| 4 | 157 | 163 | 164 | 179 |

[5 rows x 785 columns]

```
[4]: train.isnull()
```

```
[4]:
```

| | label | pixel1 | pixel2 | pixel3 | pixel4 | pixel5 | pixel6 | pixel7 | pixel8 | \ |
|-------|-------|--------|--------|--------|--------|--------|--------|--------|--------|---|
| 0 | False | False | False | False | False | False | False | False | False | |
| 1 | False | False | False | False | False | False | False | False | False | |
| 2 | False | False | False | False | False | False | False | False | False | |
| 3 | False | False | False | False | False | False | False | False | False | |
| 4 | False | False | False | False | False | False | False | False | False | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 27450 | False | False | False | False | False | False | False | False | False | |
| 27451 | False | False | False | False | False | False | False | False | False | |
| 27452 | False | False | False | False | False | False | False | False | False | |
| 27453 | False | False | False | False | False | False | False | False | False | |
| 27454 | False | False | False | False | False | False | False | False | False | |

| | pixel19 | ... | pixel775 | pixel776 | pixel777 | pixel778 | pixel779 | \ |
|-------|---------|-----|----------|----------|----------|----------|----------|---|
| 0 | False | ... | False | False | False | False | False | |
| 1 | False | ... | False | False | False | False | False | |
| 2 | False | ... | False | False | False | False | False | |
| 3 | False | ... | False | False | False | False | False | |
| 4 | False | ... | False | False | False | False | False | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 27450 | False | ... | False | False | False | False | False | |
| 27451 | False | ... | False | False | False | False | False | |
| 27452 | False | ... | False | False | False | False | False | |
| 27453 | False | ... | False | False | False | False | False | |
| 27454 | False | ... | False | False | False | False | False | |

| | pixel780 | pixel781 | pixel782 | pixel783 | pixel784 |
|-------|----------|----------|----------|----------|----------|
| 0 | False | False | False | False | False |
| 1 | False | False | False | False | False |
| 2 | False | False | False | False | False |
| 3 | False | False | False | False | False |
| 4 | False | False | False | False | False |
| ... | ... | ... | ... | ... | ... |
| 27450 | False | False | False | False | False |
| 27451 | False | False | False | False | False |
| 27452 | False | False | False | False | False |
| 27453 | False | False | False | False | False |
| 27454 | False | False | False | False | False |

[27455 rows x 785 columns]

```
[5]: train.isna().sum()
```

```
[5]: label      0
    pixel1      0
    pixel2      0
    pixel3      0
    pixel4      0
    ..
    pixel780    0
    pixel781    0
    pixel782    0
    pixel783    0
    pixel784    0
    Length: 785, dtype: int64
```

```
[6]: train_df_original = train.copy()

    # Split into training, test and validation sets
    val_index = int(train.shape[0]*0.2)

    train_df = train_df_original.iloc[val_index:]
    val_df = train_df_original.iloc[:val_index]
```

```
[7]: y = np.array(train_df['label'])
    X = np.array(train_df.drop(columns='label'))
```

```
[8]: X.shape,y.shape
```

```
[8]: ((21964, 784), (21964,))
```

```
[9]: import random
    r = random.randint(0,(21964-1))
    def show_img():
        arr = np.array(X)
        some_value = arr[r]
        some_img = some_value.reshape(28,28)
        plt.imshow(some_img, cmap="gray")
        plt.axis("off")
        plt.show()

    show_img()
    print(y[r])
```



12

```
[10]: y_train = pd.get_dummies(y)
      y_train.head(5)
```

```
[10]:
```

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 10 | ... | 15 | 16 | 17 | 18 | 19 | 20 | 21 | \ |
|---|---|---|---|---|---|---|---|---|---|----|-----|----|----|----|----|----|----|----|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0 | 1 | 0 | 0 | 0 | 0 | |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 1 | 0 | 0 | 0 | 0 | 0 | 0 | |

| | 22 | 23 | 24 |
|---|----|----|----|
| 0 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 |
| 2 | 0 | 0 | 0 |
| 3 | 0 | 0 | 0 |
| 4 | 0 | 0 | 0 |

[5 rows x 24 columns]

```
[11]: y_val = val_df['label']
      X_val = val_df.drop(columns="label",axis=1)
```

```
[12]: y_val = pd.get_dummies(y_val)
```

```
[13]: y_train.shape
```

```
[13]: (21964, 24)
```

```
[14]: X_val = pd.DataFrame(X_val).values.reshape(X_val.shape[0] ,28, 28, 1)
```

```
[15]: X_train = pd.DataFrame(X).values.reshape(X.shape[0] ,28, 28, 1)
```

```
[16]: X_train.shape,y_train.shape
```

```
[16]: ((21964, 28, 28, 1), (21964, 24))
```

```
[17]: generator = tf.keras.preprocessing.image.ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=10,  
    zoom_range=0.10,  
    width_shift_range=0.1,  
    height_shift_range=0.1,  
    shear_range=0.1,  
    horizontal_flip=False,  
    fill_mode="nearest"  
)  
  
X_train_flow = generator.flow(X_train, y_train, batch_size=32)  
  
X_val_flow = generator.flow(X_val, y_val, batch_size=32)
```

```
[18]: model = Sequential()  
  
model.add(Conv2D(filters=32, kernel_size=(3,3), activation="relu",  
    ↪input_shape=(28,28,1)))  
model.add(MaxPool2D((2,2),padding='SAME'))  
model.add(Dropout(rate=0.2))  
  
model.add(Conv2D(filters=64, kernel_size=(3,3), activation="relu",  
    ↪input_shape=(28,28,1)))  
model.add(MaxPool2D((2,2),padding='SAME'))  
model.add(Dropout(rate=0.2))  
  
model.add(Conv2D(filters=521, kernel_size=(3,3), activation="relu",  
    ↪input_shape=(28,28,1)))  
model.add(MaxPool2D((2,2),padding='SAME'))  
model.add(Dropout(rate=0.2))  
  
model.add(Flatten())
```

```

model.add(Dense(units=521, activation="relu"))
model.add(Dense(units=256, activation="relu"))
model.add(Dropout(0.5))
model.add(Dense(units=24, activation="softmax"))

model.compile(loss="categorical_crossentropy", optimizer='adam',
              metrics=["accuracy"])

```

```
[19]: model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| dropout (Dropout) | (None, 13, 13, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 11, 11, 64) | 18496 |
| max_pooling2d_1 (MaxPooling2D) | (None, 6, 6, 64) | 0 |
| dropout_1 (Dropout) | (None, 6, 6, 64) | 0 |
| conv2d_2 (Conv2D) | (None, 4, 4, 521) | 300617 |
| max_pooling2d_2 (MaxPooling2D) | (None, 2, 2, 521) | 0 |
| dropout_2 (Dropout) | (None, 2, 2, 521) | 0 |
| flatten (Flatten) | (None, 2084) | 0 |
| dense (Dense) | (None, 521) | 1086285 |
| dense_1 (Dense) | (None, 256) | 133632 |
| dropout_3 (Dropout) | (None, 256) | 0 |
| dense_2 (Dense) | (None, 24) | 6168 |

```
Total params: 1,545,518  
Trainable params: 1,545,518  
Non-trainable params: 0
```

```
[28]: dot_img_file = './SignLanguageRecognitionModel.png'  
      tf.keras.utils.plot_model(model, to_file=dot_img_file, show_shapes=True)
```

```
[28]:
```

| | | | |
|--------------|---------|---------------------|---------------------|
| conv2d_input | input: | [(None, 28, 28, 1)] | [(None, 28, 28, 1)] |
| InputLayer | output: | | |



| | | | |
|--------|---------|-------------------|--------------------|
| conv2d | input: | (None, 28, 28, 1) | (None, 26, 26, 32) |
| Conv2D | output: | | |



| | | | |
|---------------|---------|--------------------|--------------------|
| max_pooling2d | input: | (None, 26, 26, 32) | (None, 13, 13, 32) |
| MaxPooling2D | output: | | |



| | | | |
|---------|---------|--------------------|--------------------|
| dropout | input: | (None, 13, 13, 32) | (None, 13, 13, 32) |
| Dropout | output: | | |



| | | | |
|----------|---------|--------------------|--------------------|
| conv2d_1 | input: | (None, 13, 13, 32) | (None, 11, 11, 64) |
| Conv2D | output: | | |



| | | | |
|-----------------|---------|--------------------|------------------|
| max_pooling2d_1 | input: | (None, 11, 11, 64) | (None, 6, 6, 64) |
| MaxPooling2D | output: | | |



| | | | |
|-----------|---------|------------------|------------------|
| dropout_1 | input: | (None, 6, 6, 64) | (None, 6, 6, 64) |
| Dropout | output: | | |



| | | | |
|----------|---------|------------------|-------------------|
| conv2d_2 | input: | (None, 6, 6, 64) | (None, 4, 4, 512) |
| Conv2D | output: | | |



| | | | |
|-----------------|---------|-------------------|-------------------|
| max_pooling2d_2 | input: | (None, 4, 4, 512) | (None, 2, 2, 512) |
| MaxPooling2D | output: | | |



| | | | |
|-----------|---------|-------------------|-------------------|
| dropout_2 | input: | (None, 2, 2, 512) | (None, 2, 2, 512) |
| Dropout | output: | | |



| | | | |
|---------|---------|-------------------|--------------|
| flatten | input: | (None, 2, 2, 512) | (None, 2048) |
| Flatten | output: | | |



| | | | |
|-------|---------|--------------|-------------|
| dense | input: | (None, 2048) | (None, 512) |
| Dense | output: | | |



| | | | |
|---------|---------|-------------|-------------|
| dense_1 | input: | (None, 512) | (None, 256) |
| Dense | output: | | |



| | | | |
|-----------|---------|-------------|-------------|
| dropout_3 | input: | (None, 256) | (None, 256) |
| Dropout | output: | | |



| | | | |
|---------|---------|-------------|------------|
| dense_2 | input: | (None, 256) | (None, 24) |
| Dense | output: | | |


```
[20]: learning_rate_reduction = tf.keras.callbacks.ReduceLROnPlateau(  
      monitor='val_accuracy', patience = 2, verbose=1, factor=0.5, min_lr=0.00001  
      )
```

```
[21]: history = model.fit(  
      X_train_flow,  
      validation_data=X_val_flow,  
      epochs=100,  
      callbacks=[  
          tf.keras.callbacks.EarlyStopping(  
              monitor='val_loss',  
              patience=5,  
              restore_best_weights=True  
          ),  
          learning_rate_reduction  
      ])
```

Epoch 1/100

687/687 [=====] - 29s 41ms/step - loss: 2.5028 -
accuracy: 0.2132 - val_loss: 0.9107 - val_accuracy: 0.6933 - lr: 0.0010

Epoch 2/100

687/687 [=====] - 29s 42ms/step - loss: 0.7890 -
accuracy: 0.7263 - val_loss: 0.3336 - val_accuracy: 0.8907 - lr: 0.0010

Epoch 3/100

687/687 [=====] - 30s 43ms/step - loss: 0.3902 -
accuracy: 0.8669 - val_loss: 0.1397 - val_accuracy: 0.9579 - lr: 0.0010

Epoch 4/100

687/687 [=====] - 28s 41ms/step - loss: 0.2517 -
accuracy: 0.9152 - val_loss: 0.0814 - val_accuracy: 0.9745 - lr: 0.0010

Epoch 5/100

687/687 [=====] - 26s 38ms/step - loss: 0.1832 -
accuracy: 0.9403 - val_loss: 0.0561 - val_accuracy: 0.9825 - lr: 0.0010

Epoch 6/100

687/687 [=====] - 26s 38ms/step - loss: 0.1482 -
accuracy: 0.9509 - val_loss: 0.0297 - val_accuracy: 0.9922 - lr: 0.0010

Epoch 7/100

687/687 [=====] - 26s 38ms/step - loss: 0.1325 -
accuracy: 0.9576 - val_loss: 0.0240 - val_accuracy: 0.9938 - lr: 0.0010

Epoch 8/100

687/687 [=====] - 26s 38ms/step - loss: 0.1058 -
accuracy: 0.9667 - val_loss: 0.0168 - val_accuracy: 0.9951 - lr: 0.0010

Epoch 9/100

687/687 [=====] - 26s 38ms/step - loss: 0.1039 -
accuracy: 0.9677 - val_loss: 0.0112 - val_accuracy: 0.9975 - lr: 0.0010

Epoch 10/100

687/687 [=====] - 26s 38ms/step - loss: 0.0850 - accuracy: 0.9734 - val_loss: 0.0145 - val_accuracy: 0.9954 - lr: 0.0010
Epoch 11/100

687/687 [=====] - 26s 38ms/step - loss: 0.0796 - accuracy: 0.9749 - val_loss: 0.0072 - val_accuracy: 0.9980 - lr: 0.0010
Epoch 12/100

687/687 [=====] - 26s 38ms/step - loss: 0.0803 - accuracy: 0.9761 - val_loss: 0.0180 - val_accuracy: 0.9949 - lr: 0.0010
Epoch 13/100

686/687 [=====>.] - ETA: 0s - loss: 0.0736 - accuracy: 0.9783
Epoch 13: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.

687/687 [=====] - 26s 38ms/step - loss: 0.0741 - accuracy: 0.9782 - val_loss: 0.0247 - val_accuracy: 0.9925 - lr: 0.0010
Epoch 14/100

687/687 [=====] - 26s 38ms/step - loss: 0.0408 - accuracy: 0.9883 - val_loss: 0.0026 - val_accuracy: 0.9987 - lr: 5.0000e-04
Epoch 15/100

687/687 [=====] - 26s 38ms/step - loss: 0.0381 - accuracy: 0.9888 - val_loss: 0.0036 - val_accuracy: 0.9987 - lr: 5.0000e-04
Epoch 16/100

687/687 [=====] - 26s 38ms/step - loss: 0.0348 - accuracy: 0.9886 - val_loss: 0.0048 - val_accuracy: 0.9989 - lr: 5.0000e-04
Epoch 17/100

687/687 [=====] - 26s 38ms/step - loss: 0.0351 - accuracy: 0.9897 - val_loss: 0.0019 - val_accuracy: 0.9995 - lr: 5.0000e-04
Epoch 18/100

687/687 [=====] - 26s 38ms/step - loss: 0.0349 - accuracy: 0.9900 - val_loss: 0.0046 - val_accuracy: 0.9989 - lr: 5.0000e-04
Epoch 19/100

687/687 [=====] - 31s 45ms/step - loss: 0.0310 - accuracy: 0.9911 - val_loss: 0.0015 - val_accuracy: 0.9998 - lr: 5.0000e-04
Epoch 20/100

687/687 [=====] - 27s 39ms/step - loss: 0.0265 - accuracy: 0.9919 - val_loss: 0.0028 - val_accuracy: 0.9996 - lr: 5.0000e-04
Epoch 21/100

687/687 [=====] - ETA: 0s - loss: 0.0268 - accuracy: 0.9920
Epoch 21: ReduceLROnPlateau reducing learning rate to 0.0002500000118743628.

687/687 [=====] - 27s 39ms/step - loss: 0.0268 - accuracy: 0.9920 - val_loss: 0.0092 - val_accuracy: 0.9971 - lr: 5.0000e-04
Epoch 22/100

687/687 [=====] - 27s 39ms/step - loss: 0.0190 - accuracy: 0.9942 - val_loss: 0.0010 - val_accuracy: 0.9998 - lr: 2.5000e-04
Epoch 23/100

687/687 [=====] - ETA: 0s - loss: 0.0147 - accuracy: 0.9953
Epoch 23: ReduceLROnPlateau reducing learning rate to 0.0001250000059371814.

687/687 [=====] - 27s 39ms/step - loss: 0.0147 -
accuracy: 0.9953 - val_loss: 0.0034 - val_accuracy: 0.9987 - lr: 2.5000e-04
Epoch 24/100

687/687 [=====] - 27s 39ms/step - loss: 0.0141 -
accuracy: 0.9958 - val_loss: 6.4741e-04 - val_accuracy: 0.9998 - lr: 1.2500e-04
Epoch 25/100

686/687 [=====>.] - ETA: 0s - loss: 0.0132 - accuracy:
0.9959
Epoch 25: ReduceLROnPlateau reducing learning rate to 6.25000029685907e-05.

687/687 [=====] - 27s 39ms/step - loss: 0.0132 -
accuracy: 0.9959 - val_loss: 0.0011 - val_accuracy: 0.9995 - lr: 1.2500e-04
Epoch 26/100

687/687 [=====] - 27s 39ms/step - loss: 0.0115 -
accuracy: 0.9969 - val_loss: 2.8248e-04 - val_accuracy: 1.0000 - lr: 6.2500e-05
Epoch 27/100

687/687 [=====] - 26s 37ms/step - loss: 0.0114 -
accuracy: 0.9965 - val_loss: 4.3466e-04 - val_accuracy: 0.9998 - lr: 6.2500e-05
Epoch 28/100

687/687 [=====] - ETA: 0s - loss: 0.0098 - accuracy:
0.9971
Epoch 28: ReduceLROnPlateau reducing learning rate to 3.125000148429535e-05.

687/687 [=====] - 26s 37ms/step - loss: 0.0098 -
accuracy: 0.9971 - val_loss: 3.3614e-04 - val_accuracy: 0.9998 - lr: 6.2500e-05
Epoch 29/100

687/687 [=====] - 27s 39ms/step - loss: 0.0123 -
accuracy: 0.9963 - val_loss: 1.6902e-04 - val_accuracy: 1.0000 - lr: 3.1250e-05
Epoch 30/100

687/687 [=====] - ETA: 0s - loss: 0.0099 - accuracy:
0.9973
Epoch 30: ReduceLROnPlateau reducing learning rate to 1.5625000742147677e-05.

687/687 [=====] - 27s 39ms/step - loss: 0.0099 -
accuracy: 0.9973 - val_loss: 2.1925e-04 - val_accuracy: 0.9998 - lr: 3.1250e-05
Epoch 31/100

687/687 [=====] - 27s 39ms/step - loss: 0.0103 -
accuracy: 0.9965 - val_loss: 1.2029e-04 - val_accuracy: 1.0000 - lr: 1.5625e-05
Epoch 32/100

686/687 [=====>.] - ETA: 0s - loss: 0.0089 - accuracy:
0.9974
Epoch 32: ReduceLROnPlateau reducing learning rate to 1e-05.

687/687 [=====] - 27s 39ms/step - loss: 0.0089 -
accuracy: 0.9975 - val_loss: 8.1509e-05 - val_accuracy: 1.0000 - lr: 1.5625e-05
Epoch 33/100

687/687 [=====] - 27s 39ms/step - loss: 0.0087 -
accuracy: 0.9974 - val_loss: 1.1326e-04 - val_accuracy: 1.0000 - lr: 1.0000e-05
Epoch 34/100

687/687 [=====] - 27s 39ms/step - loss: 0.0079 -
accuracy: 0.9971 - val_loss: 1.4355e-04 - val_accuracy: 1.0000 - lr: 1.0000e-05
Epoch 35/100

```

687/687 [=====] - 28s 41ms/step - loss: 0.0097 -
accuracy: 0.9970 - val_loss: 0.0013 - val_accuracy: 0.9995 - lr: 1.0000e-05
Epoch 36/100
687/687 [=====] - 27s 40ms/step - loss: 0.0085 -
accuracy: 0.9976 - val_loss: 7.3692e-04 - val_accuracy: 0.9996 - lr: 1.0000e-05
Epoch 37/100
687/687 [=====] - 27s 39ms/step - loss: 0.0091 -
accuracy: 0.9973 - val_loss: 2.8997e-04 - val_accuracy: 1.0000 - lr: 1.0000e-05

```


```

[22]: fig, axes = plt.subplots(2, 1, figsize=(15, 10))
      ax = axes.flat

      pd.DataFrame(history.history)[['accuracy', 'val_accuracy']].plot(ax=ax[0])
      ax[0].set_title("Accuracy", fontsize = 15)
      ax[0].set_ylim(0,1.1)

      pd.DataFrame(history.history)[['loss', 'val_loss']].plot(ax=ax[1])
      ax[1].set_title("Loss", fontsize = 15)
      plt.show()

```



```

[23]: y_test = np.array(test['label'])
      X_test = np.array(test.drop(columns='label'))

```

```

y_test = pd.get_dummies(y_test)
X_test = pd.DataFrame(X_test).values.reshape(X_test.shape[0] ,28, 28, 1)

# X_test_flow = generator.flow(X_test, y_test, batch_size=32)
# X_test.shape,X_train.shape

y_test = pd.get_dummies(y_test)

```

```
[24]: from sklearn.metrics import classification_report
```

```

# predictions
pred = model.predict(X_test)

y_pred = np.argmax(pred,axis=1)
y_test = np.argmax(y_test.values,axis=1)

```

```
[25]: acc = accuracy_score(y_test,y_pred)
```

```

# # Display the results
print(f'## {acc*100:.2f}% accuracy on the test set')

```

99.51% accuracy on the test set

```
[32]: fer_json = model.to_json()
with open("./SavedModel/SignLanguageRecognitionModel.json", "w") as json_file:
    json_file.write(fer_json)
model.save('./SavedModel/SignLanguageRecognitionModel_tf',save_format='tf')
model.save("./SavedModel/SignLanguageRecognitionModel.h5")
model.save_weights("./SavedModel/SignLanguageRecognitionModel.h5")

```

INFO:tensorflow:Assets written to:
./SavedModel/SignLanguageRecognitionModel_tf/assets